

AnalyseHeath

January 15, 2023

1 Logistic Regression for health

1.1 Comprehension of caracteristiques

It will use real-world data that contains detailed nutritional information about foods for people with diabetes. The goal is to determine whether a diabetic patient should choose more often, less often, or in moderation for a specific food item based on the nutritional information in the dataset.

1.2 Principale objectif and mission

I would like to work for a hospital as a data scientist and I have no data in this field. So I took the data from the training because I am also a student for the moment. The goal is to determine whether a diabetic patient should choose more often, less often, or in moderation for a specific food item based on the nutritional information in the dataset. My analysis objective is mainly to retrieve data related to diabetes and to analyze : * Train and fine-tune logistic regression models * Interpret trained logistic regression models * Evaluate trained logistic regression models * The explanatory factors of the degradation and concentration of the diabetes rate in the blood * Make recommandation for next setps

1.2.1 Importation for librairies and preparation

```
[34]: import pandas as pd
import numpy as np
from sklearn.preprocessing import OneHotEncoder, LabelEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split, learning_curve
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import classification_report, accuracy_score,
    ↪confusion_matrix, precision_recall_fscore_support, precision_score,
    ↪recall_score
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[35]: dataset_url = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.
    ↪cloud/IBM-ML241EN-SkillsNetwork/labs/datasets/food_items.csv"
food_df = pd.read_csv(dataset_url)
```

Observations of variables and data values

```
[36]: food_df.head(10)
```

```
[36]:
```

	Calories	Total Fat	Saturated Fat	Monounsaturated Fat	\
0	149.0	0	0.0	0.0	
1	123.0	0	0.0	0.0	
2	150.0	0	0.0	0.0	
3	110.0	0	0.0	0.0	
4	143.0	0	0.0	0.0	
5	110.0	0	0.0	0.0	
6	142.0	0	0.0	0.0	
7	102.0	0	0.0	0.0	
8	145.0	0	0.0	0.0	
9	171.0	0	0.0	0.0	

	Polyunsaturated Fat	Trans Fat	Cholesterol	Sodium	Total Carbohydrate	\
0	0.0	0.0	0	9.0	9.8	
1	0.0	0.0	0	5.0	6.6	
2	0.0	0.0	0	4.0	11.4	
3	0.0	0.0	0	6.0	7.0	
4	0.0	0.0	0	7.0	13.1	
5	0.0	0.0	0	6.0	7.0	
6	0.0	0.0	0	12.0	10.6	
7	0.0	0.0	0	13.0	5.0	
8	0.0	0.0	0	17.0	11.0	
9	0.0	0.0	0	8.0	13.7	

	Dietary Fiber	Sugars	Sugar Alcohol	Protein	Vitamin A	Vitamin C	\
0	0.0	0.0	0	1.3	0	0	
1	0.0	0.0	0	0.8	0	0	
2	0.0	0.0	0	1.3	0	0	
3	0.0	0.0	0	0.8	0	0	
4	0.0	0.0	0	1.0	0	0	
5	0.0	0.0	0	0.8	0	0	
6	0.0	0.0	0	1.2	0	0	
7	0.0	0.0	0	0.7	0	0	
8	0.0	0.0	0	1.2	0	0	
9	0.0	0.0	0	2.5	0	0	

	Calcium	Iron	class
0	0	0	'In Moderation'
1	0	0	'In Moderation'
2	0	0	'In Moderation'
3	0	0	'In Moderation'
4	0	0	'In Moderation'
5	0	0	'In Moderation'
6	0	0	'In Moderation'
7	0	0	'In Moderation'

8	0	0	'In Moderation'
9	0	0	'In Moderation'

1.2.2 Data mean

- Calories Calories of patients;
- Total Fat Total Fat;
- Saturated Fat Saturated ;
- Monounsaturated Fat Monounsaturated ;
- Polyunsaturated Fat Polyunsaturated ;
- Trans Fat Trans Fat ;
- Cholesterol Presence of cholesterol ;
- Sodium Sodium quantity ;
- Total Carbohydrate Carbonnate quantity ;
- Dietary Fiber fievre ;
- Sugars Sugars quantity ;
- Sugar Alcohol Presence of Alcool ;
- Protein Proteine ;
- Vitamin A Vitamin ;
- Vitamin C Vitamin ;
- Calcium Calorie ;
- Iron Calorie ;
- class predictor and explain variable

1.2.3 We explain categories of class with this variables

- Calories Calories of patients;
- Total Fat Total Fat;
- Saturated Fat Saturated ;
- Monounsaturated Fat Monounsaturated ;
- Polyunsaturated Fat Polyunsaturated ;
- Trans Fat Trans Fat ;
- Cholesterol Presence of cholesterol ;
- Sodium Sodium quantity ;
- Total Carbohydrate Carbonnate quantity ;
- Dietary Fiber fievre ;
- Sugars Sugars quantity ;
- Sugar Alcohol Presence of Alcool ;
- Protein Proteine ;
- Vitamin A Vitamin ;
- Vitamin C Vitamin ;
- Calcium Calorie ;
- Iron Calorie ;

1.2.4 Type of data

```
[24]: food_df.dtypes
```

```
[24]: Calories                float64
      Total Fat              int64
      Saturated Fat         float64
      Monounsaturated Fat   float64
      Polyunsaturated Fat   float64
      Trans Fat             float64
      Cholesterol           int64
      Sodium               float64
      Total Carbohydrate    float64
      Dietary Fiber         float64
      Sugars                float64
      Sugar Alcohol         int64
      Protein              float64
      Vitamin A             int64
      Vitamin C             int64
      Calcium              int64
      Iron                 int64
      class                 object
      dtype: object
```

1.2.5 Description statistics

```
[37]: food_df.describe()
```

```
[37]:
```

	Calories	Total Fat	Saturated Fat	Monounsaturated Fat	\
count	13260.000000	13260.000000	13260.000000	13260.000000	
mean	133.861086	4.475264	1.450617	0.338069	
std	94.227650	5.386340	2.410318	1.345852	
min	0.000000	0.000000	0.000000	0.000000	
25%	70.000000	0.000000	0.000000	0.000000	
50%	120.000000	3.000000	0.500000	0.000000	
75%	180.000000	7.000000	2.000000	0.000000	
max	2210.000000	43.000000	22.000000	40.000000	

	Polyunsaturated Fat	Trans Fat	Cholesterol	Sodium	\
count	13260.000000	13260.000000	13260.000000	13260.000000	
mean	0.254660	0.047459	8.857692	241.867142	
std	2.230586	0.321402	20.976530	272.284363	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	40.000000	
50%	0.000000	0.000000	0.000000	135.000000	
75%	0.000000	0.000000	10.000000	370.000000	
max	235.000000	11.000000	450.000000	2431.000000	

	Total Carbohydrate	Dietary Fiber	Sugars	Sugar Alcohol	\
count	13260.000000	13260.000000	13260.000000	13260.000000	
mean	18.232020	1.602971	6.645234	0.117949	
std	14.786316	3.363879	8.328465	1.121529	
min	0.000000	0.000000	0.000000	0.000000	
25%	5.000000	0.000000	0.000000	0.000000	
50%	17.000000	1.000000	3.000000	0.000000	
75%	27.000000	2.000000	11.000000	0.000000	
max	270.000000	305.000000	115.000000	31.000000	

	Protein	Vitamin A	Vitamin C	Calcium	Iron
count	13260.000000	13260.000000	13260.000000	13260.000000	13260.000000
mean	4.661333	6.287632	6.741855	5.175264	5.235671
std	5.611143	18.374191	23.785100	8.779637	9.119459
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000
50%	3.000000	0.000000	0.000000	2.000000	2.000000
75%	7.000000	6.000000	2.000000	6.000000	8.000000
max	70.000000	622.000000	1000.000000	110.000000	170.000000

1.3 Variante models and classification choice

After presentation we are any missing data. As we can see from the above output, this dataset contains 17 nutrient categories about each food item. These categories include Calories, Total Fat, Protein, Sugar, etc., and are listed as numeric variables. As such, we only need to scale them for training our logistic regression model so that we can compare our feature coefficients directly.

We have three labels meaning our logistic regression model will be multinomial with three classes.

A multinomial logistic regression is a generalized logistic regression model which generates a probability distribution over all classes, based on the logits or exponentiated log-odds calculated for each class (usually more than two). We can try : * logistic regression * Decisions Tree * SVMs * KNNs * Ensemble learning like Random Forest

But we start with logistic regression. Also note that a multinomial logistic regression model is different from the **one-vs-rest** binary logistic regression. For **one-vs-rest** schema, you need to train an independent classifier for each class. For example, you need a **More Often** classifier to differentiate a food item between **More Often** and **Not More Often** (or, **In Moderation** and **Less Often**).


1.3.1 Predictor

```
[38]: food_df['class'].value_counts(normalize=True)
```

```
[38]: 'In Moderation'    0.501433
      'Less Often'      0.423906
      'More Often'      0.074661
      Name: class, dtype: float64
```

```
[39]: food_df['class'].value_counts().plot.barh(color=['yellow', 'red', 'green'])
```

```
[39]: <AxesSubplot:>
```



We can see on the graph above, this data set has three classes: “In moderation”, “Less often” and “More often”. All three labels are unbalanced. For diabetic patients, most foods fall into the “In moderation” and “Less often” categories. This makes managing the diabetic diet very difficult. Therefore, we could build a machine learning model to help patients choose their foods.

We have three labels meaning our logistic regression model will be multinomial with three classes.

1.4 Feature engennering

Fortunately, all feature columns are numeric so we just need to scale them. Here we use the `MinMaxScaler` provided by `sklearn` for scaling.

```
[40]: #MinMaxScaler object  
scaler = MinMaxScaler()
```

```
[41]: #Application with data after identification  
X_raw = food_df.iloc[:, :-1]  
y_raw = food_df.iloc[:, -1:]  
X = scaler.fit_transform(X_raw)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-  
packages/sklearn/preprocessing/data.py:323: DataConversionWarning: Data with  
input dtype int64, float64 were all converted to float64 by MinMaxScaler.  
    return self.partial_fit(X, y)
```

For the target variable `y`, let's use the `LabelEncoder` provided by `sklearn` to encode its three class values.

```
[42]: # LabelEncoder object  
label_encoder = LabelEncoder()
```

```
[43]: #Applplication of data  
y = label_encoder.fit_transform(y_raw.values.ravel())
```

The encoded target variable will only contain values 0=In Moderation, 1=Less Often, 2=More Often.

```
[44]: #Note of variable y 0  
np.unique(y, return_counts=True)
```

```
[44]: (array([0, 1, 2]), array([6649, 5621, 990]))
```

1.5 Preparation of modelisation

```
[45]: # First, let's split the training and testing dataset  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
    ↪stratify=y, random_state = 10)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-  
packages/sklearn/model_selection/_split.py:1609: DeprecationWarning: `np.int` is  
a deprecated alias for the builtin `int`. To silence this warning, use `int` by  
itself. Doing this will not modify any behavior and is safe. When replacing  
`np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the  
precision. If you wish to review your current use, check the release note link  
for additional information.
```

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
    return floored.astype(np.int)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-  
packages/sklearn/model_selection/_split.py:1609: DeprecationWarning: `np.int` is  
a deprecated alias for the builtin `int`. To silence this warning, use `int` by  
itself. Doing this will not modify any behavior and is safe. When replacing  
`np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the  
precision. If you wish to review your current use, check the release note link  
for additional information.
```

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
    return floored.astype(np.int)
```

```
[46]: # L2 penalty
penalty= 'l2'
# Type of classification problem is multinomial
multi_class = 'multinomial'
# Use lbfgs for L2 penalty and multinomial classes
solver = 'lbfgs'
# Max iteration = 1000
max_iter = 1000
```

```
[47]: # Define a logistic regression model with above arguments
l2_model = LogisticRegression(random_state=10, penalty=penalty,
    ↪multi_class=multi_class, solver=solver, max_iter=max_iter)
```

```
[48]: # Creation of mofel
l2_model.fit(X_train, y_train)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-
packages/sklearn/utils/fixes.py:357: DeprecationWarning: distutils Version
classes are deprecated. Use packaging.version instead.
    if _joblib.__version__ >= LooseVersion('0.12'):
```

```
[48]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=1000, multi_class='multinomial',
    n_jobs=None, penalty='l2', random_state=10, solver='lbfgs',
    tol=0.0001, verbose=0, warm_start=False)
```

```
[49]: l2_preds = l2_model.predict(X_test)
```

```
[50]: l2_preds
```

```
[50]: array([1, 0, 2, ..., 0, 0, 2])
```

1.6 Evaluation

```
[51]: # Function of evaluation
def evaluate_metrics(yt, yp):
    results_pos = {}
    results_pos['accuracy'] = accuracy_score(yt, yp)
    precision, recall, f_beta, _ = precision_recall_fscore_support(yt, yp)
    results_pos['recall'] = recall
    results_pos['precision'] = precision
    results_pos['f1score'] = f_beta
    return results_pos
```

```
[52]: evaluate_metrics(y_test, l2_preds)
```



```
[52]: {'accuracy': 0.7669683257918553,
      'recall': array([0.87067669, 0.73220641, 0.26767677]),
      'precision': array([0.72194514, 0.83047427, 0.92982456]),
      'f1score': array([0.78936605, 0.77825059, 0.41568627])}
```

As we can see from the above evaluation results, the logistic regression model has relatively good performance on this multinomial classification task. The overall accuracy is around 0.76 and the f1score is around 0.7. Note that for **recall**, **precision**, and **f1score**, we output the values for each class to see how the model performs on an individual class. And, we can see from the results, the recall for **class=2** (More often) is not very good. This is actually a common problem called imbalanced classification challenge. We will introduce solution to this problem later in this course.

1.7 We can try again for next performance

```
[53]: # L1 penalty
      penalty= 'l1'
      # Our classification problem is multinomial
      multi_class = 'multinomial'
      # Use saga for L1 penalty and multinomial classes
      solver = 'saga'
      # Max iteration = 1000
      max_iter = 1000
```

```
[54]: # Define a logistic regression model with above arguments
      l2_model = LogisticRegression(random_state=120, penalty=penalty,
      ↪ multi_class=multi_class, solver=solver, max_iter = 1000)
```

```
[55]: l2_model.fit(X_train, y_train)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.7/site-
packages/sklearn/utils/fixes.py:357: DeprecationWarning: distutils Version
classes are deprecated. Use packaging.version instead.
    if _joblib.__version__ >= LooseVersion('0.12'):
```

```
[55]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
      intercept_scaling=1, max_iter=1000, multi_class='multinomial',
      n_jobs=None, penalty='l1', random_state=120, solver='saga',
      tol=0.0001, verbose=0, warm_start=False)
```

```
[56]: l2_preds = l2_model.predict(X_test)
```

```
[57]: l2_preds
```

```
[57]: array([1, 0, 2, ..., 0, 0, 2])
```

1.8 Application for interpretation

```
[58]: odd_ratios = l2_model.predict_proba(X_test[:, :])[0]
      odd_ratios
```


```
[58]: array([1.78052047e-01, 8.21928521e-01, 1.94312960e-05])
```

```
[59]: evaluate_metrics(y_test, l2_preds)
```

```
[59]: {'accuracy': 0.8046757164404224,
      'recall': array([0.85639098, 0.7366548 , 0.84343434]),
      'precision': array([0.78174331, 0.8313253 , 0.83919598]),
      'f1score': array([0.81736634, 0.78113208, 0.84130982])}
```

1.9 We are amelioration

```
[60]: ## Make confusion matrix
      cf = confusion_matrix(y_test, l2_preds)
      plt.figure(figsize=(16, 12))
      ax = sns.heatmap(cf, annot=True, fmt="d", xticklabels=["In Moderation", "Less_
      ↪Often", "More Often"],
                      yticklabels=["In Moderation", "Less Often", "More Often"])
      ax.set(title="Confusion Matrix");
```



1.9.1 Interpret logistic regression models

One way to interpret logistic regression models is by analyzing feature coefficients. Although it may not be as effective as the regular linear regression models because the logistic regression model has a sigmoid function, we can still get a sense for the importance or impact of each feature.

1.10 Presentation

We find that, unhealthy nutrients such as saturated fat, sugars, cholesterol, total fat, etc. have high positive coefficients. Foods containing unhealthy nutrients will have higher coefficients and will be more likely to be classified as “Less often”.

```
[61]: ## Coefficients
      l2_model.coef_
```

```
[61]: array([[ 9.17517965,  0.          ,  0.          ,  5.01160431,
              0.          , -3.47379064,  0.          ,  0.65310136,
              0.          , 26.97955638,  0.          ,  0.          ,
              4.69404394,  0.21378175,  0.          ,  0.756085  ]])
```

```

    0.      ],
[  0.      ,  4.27153943, 24.0845023 ,  0.      ,
  0.      ,  0.      ,  6.76348946,  0.      ,
 1.39732665,  0.      , 15.06202754,  4.41802313,
 0.      , -2.19329857,  0.      , -1.04899691,
 0.      ],
[-118.87271714, -29.35983135,  0.      ,  0.      ,
 0.      ,  0.      ,  0.      , -1.41583778,
-44.43511175,  0.      ,  0.      ,  0.      ,
 0.      ,  0.      ,  0.      ,  0.      ,
 0.      ]])

```

```

[63]: # Extract and sort feature coefficients
def get_feature_coefs(regression_model, label_index, columns):
    coef_dict = {}
    for coef, feat in zip(regression_model.coef_[label_index, :], columns):
        if abs(coef) >= 0.01:
            coef_dict[feat] = coef
    # Sort coefficients
    coef_dict = {k: v for k, v in sorted(coef_dict.items(), key=lambda item:
    ↪item[1])}
    return coef_dict

# Generate bar colors based on if value is negative or positive
def get_bar_colors(values):
    color_vals = []
    for val in values:
        if val <= 0:
            color_vals.append('r')
        else:
            color_vals.append('g')
    return color_vals

# Visualize coefficients
def visualize_coefs(coef_dict):
    features = list(coef_dict.keys())
    values = list(coef_dict.values())
    y_pos = np.arange(len(features))
    color_vals = get_bar_colors(values)
    plt.rcParams()
    fig, ax = plt.subplots()
    ax.barh(y_pos, values, align='center', color=color_vals)
    ax.set_yticks(y_pos)
    ax.set_yticklabels(features)
    # labels read top-to-bottom
    ax.invert_yaxis()
    ax.set_xlabel('Feature Coefficients')

```

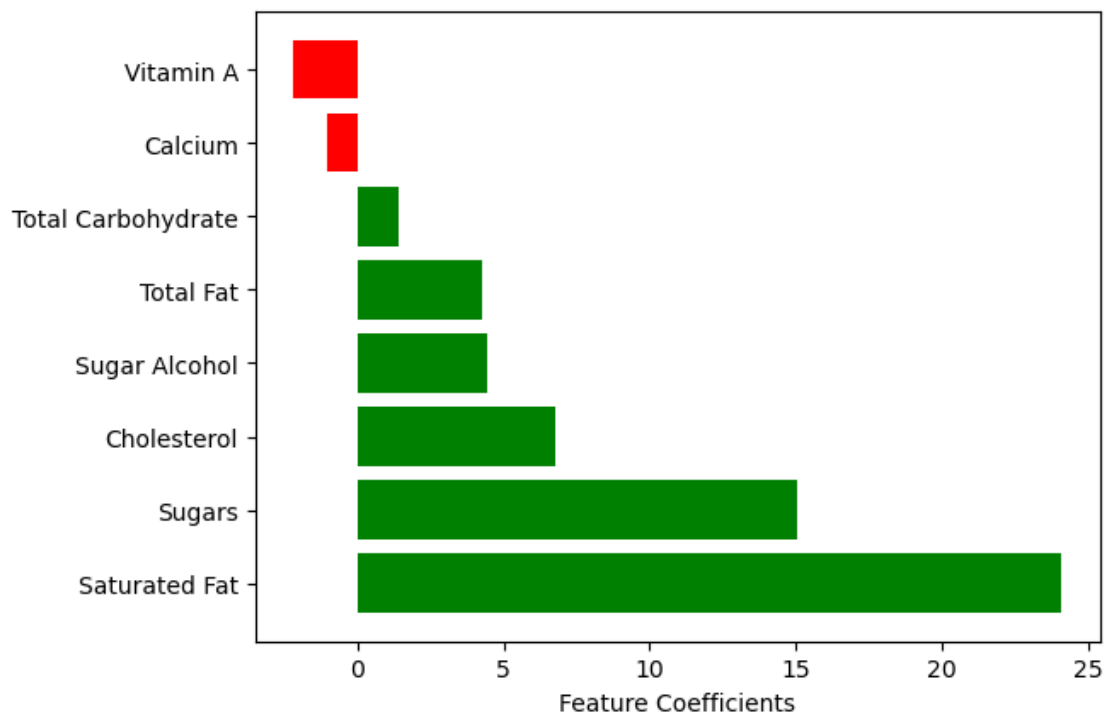
```
ax.set_title('')  
plt.show()
```

```
[68]: feature_cols = list(food_df.iloc[:, :-1].columns)  
feature_cols
```

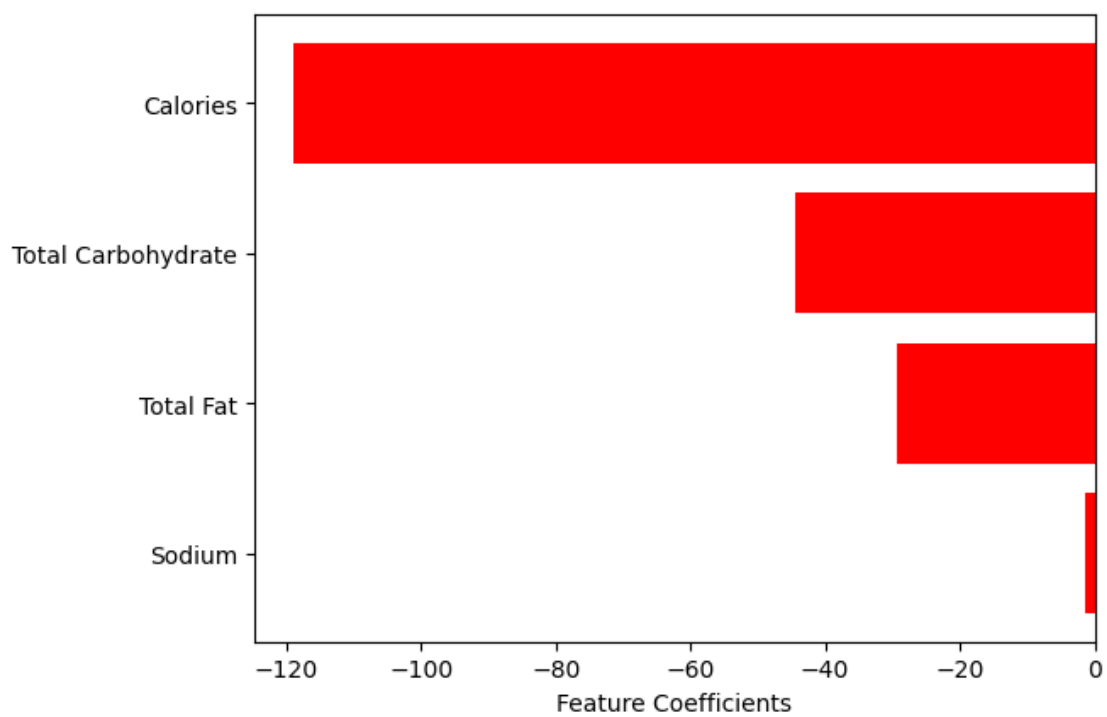
```
[68]: ['Calories',  
      'Total Fat',  
      'Saturated Fat',  
      'Monounsaturated Fat',  
      'Polyunsaturated Fat',  
      'Trans Fat',  
      'Cholesterol',  
      'Sodium',  
      'Total Carbohydrate',  
      'Dietary Fiber',  
      'Sugars',  
      'Sugar Alcohol',  
      'Protein',  
      'Vitamin A',  
      'Vitamin C',  
      'Calcium',  
      'Iron']
```

```
[71]: # Get the coefficients for Class 1, Less Often  
coef_dict = get_feature_coefs(l2_model, 1, feature_cols)
```

```
[72]: visualize_coefs(coef_dict)
```



```
[70]: # Coefficients for Class 2
coef_dict = get_feature_coefs(l2_model, 2, feature_cols)
visualize_coefs(coef_dict)
```



1.11 Conclusion

We find that, unhealthy nutrients such as saturated fat, sugars, cholesterol, total fat, etc. have high positive coefficients. Foods containing unhealthy nutrients will have higher coefficients and will be more likely to be classified as “Less often”.

For better explainability, we will try other models such as decision trees and K-nearest neighbors because overall the diabetes decision seems to be based on variables that have some relationship between them. The main weakness of this analysis is the fact that we have unbalanced classes and it will be necessary to think about the resampling in order to better apply the logistic regression

[]: